

Text-RSSR: Text-Conditioned Remote Sensing Super-Resolution with Cross-Modal Learning

Hao Yang, Man-On Pun, *Senior Member, IEEE*, and Xianping Ma

Abstract—Super-resolution of remote sensing images seeks to reconstruct high-resolution content from low-resolution inputs, thereby providing critical detail for tasks such as environmental monitoring, urban mapping, and land-cover analysis. While deep learning-based methods have recently advanced super-resolution performance, they remain limited in effectively capturing semantic context. Most existing approaches rely solely on visual information, which often leads to over-smoothed textures, inadequate structural detail, and poor generalization across diverse landscapes. Moreover, current decoder designs typically struggle to reconstruct semantically faithful content, particularly in complex remote sensing scenarios. To address these challenges, this paper proposes Text-RSSR, a text-conditioned remote sensing super-resolution framework that introduces semantic guidance through cross-modal learning. By leveraging text captions and CLIP embeddings, the model incorporates high-level contextual priors to iteratively refine visual features and enhance reconstruction quality. This integration of linguistic and visual information enables the network to achieve a stronger balance between fine detail restoration and semantic consistency. Comprehensive experiments on benchmark remote sensing datasets demonstrate that Text-RSSR outperforms state-of-the-art methods in both quantitative accuracy and perceptual quality, highlighting the effectiveness of cross-modal conditioning for advancing remote sensing super-resolution. The implementation will be made publicly available at <https://github.com/haoyangofficial/textrssr>.

Index Terms—Remote sensing image, super-resolution, multimodal learning

I. RELATED WORKS

A. RSSR Methods

RSSR aims to reconstruct high-resolution (HR) images from low-resolution (LR) observations, thereby enhancing spatial details crucial for tasks such as land cover classification, urban monitoring, and environmental analysis. Classical RSSR approaches primarily rely on signal processing or statistical modeling. Interpolation-based techniques, such as bicubic interpolation [1], provide simple baselines but often result in blurry reconstructions due to their inability to recover high-frequency details. Sparse representation methods, such as Yang *et al.*'s sparse coding-based SR [2], improve over interpolation by learning dictionaries of high- and low-resolution patches, but they are still limited in capturing the complex spatial and spectral structures inherent in remote sensing images.

This work was supported in part by the National Natural Science Foundation of China. (*Corresponding authors: Man-On Pun and Xianping Ma*)

Hao Yang and Man-On Pun are with the School of Science and Engineering, The Chinese University of Hong Kong, Shenzhen, Shenzhen 518172, China (emails: haoyang.official@gmail.com; SimonPun@cuhk.edu.cn).

Xianping Ma is with the Faculty of Geosciences and Environmental Engineering, Southwest Jiaotong University, Chengdu 611756, China (email: ma.xianping125@gmail.com).

In recent years, deep learning (DL)-based methods have achieved remarkable progress in RSSR by learning non-linear mappings from LR to HR images. Convolutional neural networks (CNNs), such as C. Dong *et al.*'s SRCNN [3] and B. Lim *et al.*'s EDSR [4], pioneered single-image RSSR by directly learning end-to-end super-resolution functions. More recently, state-space or Mamba-based architectures have emerged, such as Y. Xiao *et al.*'s FMSR [5] and T. Wu *et al.*'s EMAN [6], which model long-range dependencies efficiently. Generative adversarial networks (GANs), including C. Ledig *et al.*'s SRGAN [7] and S. Lei *et al.*'s CDGAN [8], have been widely adopted to generate perceptually realistic HR images. Diffusion-based models, such as H. Li *et al.*'s SRDiff [9] and B. Kavar *et al.*'s DDRM [10], leverage probabilistic generative frameworks to progressively refine image quality. Despite these advances, existing DL-based methods still face significant limitations: they primarily focus on unimodal visual information and rarely leverage auxiliary modalities, particularly textual descriptions, which can provide high-level semantic context about the scene.

To address this limitation, this study introduces Text-Conditioned Remote Sensing Super-Resolution (Text-RSSR), a multimodal framework designed to incorporate textual information into the RSSR process. Text-RSSR leverages vision-language models (VLMs) to generate captions from HR images, which are then encoded into embeddings using CLIP [11]. During super-resolution, the LR image is jointly processed by CNN and CLIP image encoders, while textual embeddings provide additional semantic cues. A text-guided iterative alignment module progressively refines CNN visual embeddings with multimodal guidance before decoding them into the final HR image. By integrating textual information, Text-RSSR aims to overcome the semantic gap of prior unimodal methods, offering a novel pathway toward semantically enriched and context-aware RSSR.

B. Multimodal Learning in RSSR

Multimodal Learning in RSSR refers to approaches that leverage complementary data types—such as SAR (Synthetic Aperture Radar), optical imagery, hyperspectral and multispectral images, or even textual/contextual information—to enhance the spatial and/or spectral fidelity of reconstructed high-resolution images.

ML-based RSSR methods can broadly be divided into cross-modal fusion methods and joint embedding or contrastive learning methods. Cross-modal fusion approaches seek to integrate complementary information from multiple sensors

to enhance reconstruction quality. For example, the Multi-Resolution Collaborative Fusion of SAR, Multispectral and Hyperspectral Images method combines SAR, multispectral, and hyperspectral imagery through deep fusion strategies to preserve spatial, spectral, and geometrical fidelity in reconstructed outputs, demonstrating effective cross-modal integration in remote sensing super-resolution contexts [12]. Another example involves combining low-resolution hyperspectral with high-resolution multispectral data through transformer-based fusion or retractable spatial-spectral modeling, yielding better spectral and spatial details [13]. In contrast, joint embedding and contrastive learning methods focus on learning a shared latent representation across modalities. One representative example is the multimodal contrastive learning framework for remote sensing tasks, which employs a dual-encoder architecture pre-trained on approximately one million paired Sentinel-1 (radar) and Sentinel-2 (optical) image pairs. This framework aligns features across modalities via contrastive learning and shows improved performance in downstream tasks like segmentation and land-cover mapping [14]. Another compelling approach is the Contrastive Radar-Optical Masked Autoencoders, which fuses cross-modal contrastive learning with masked autoencoding to learn rich multimodal representations from spatially aligned SAR and multispectral optical imagery [15]. Despite the promise of these fusion and embedding-based RSSR approaches, they tend to rely heavily on additional visual modalities—such as SAR, hyperspectral, or multispectral data—while largely neglecting the potential of textual information, including captions and semantic descriptions, that could provide high-level guidance for more semantically aware super-resolution.

To address this gap, this study introduces Text-RSSR—a Text-Conditioned Remote Sensing Super-Resolution Network. Text-RSSR synergizes image data with textual semantics by first generating captions for high-resolution images and encoding them via CLIP embeddings. These text embeddings, along with CNN and CLIP image embeddings, undergo an iterative Text-Guided Feature Aligning process within a Super-Resolution (SR) module. The framework effectively fuses visual and textual cues through successive refinement steps, culminating in a CNN-based image decoder that produces enhanced high-resolution outputs. By incorporating textual semantics, Text-RSSR remedies the underutilization of language in multimodal RSSR and offers more semantically informed reconstruction. While no prior RSSR method explicitly employs text conditioning, related work in natural image super-resolution demonstrates the viability of CLIP-based text guidance, as in the recent Text-Guided Explorable Image Super-Resolution framework [16]. Text-RSSR thus represents a novel extension of multimodal learning to remote sensing super-resolution by integrating visual and linguistic modalities for the first time.

II. PROPOSED METHOD

A. Overview

This study introduces a Text-Conditioned Remote Sensing Super-Resolution Network (Text-RSSR) designed to improve

remote sensing super-resolution performance through the iterative use of text captions and CLIP embeddings. As shown in Figure 1, the Text-RSSR framework comprises four main steps: Data Preparation, Multimodal Encoding, Text-Guided Iterative Super-Resolution, and Image Decoding.

In the Data Preparation step, a VLM generates a text caption from the high-resolution (HR) image, capturing its detailed content. During Multimodal Encoding, the low-resolution (LR) image is processed by two separate encoders—a CNN Image Encoder and a CLIP Image Encoder—to produce CNN Visual Embeddings and CLIP Visual Embeddings, respectively. Simultaneously, the text caption is processed through a CLIP Text Encoder to generate Text Embeddings. In the CLIP-Guided Iterative Super-Resolution step, the CNN Visual Embeddings are iteratively aligned and refined. In each iteration, all embeddings are input into a Text-Guided Feature Aligning step, followed by processing through a Super-Resolution (SR) Unit. After the final iteration, the refined CNN Visual Embeddings go through one last Text-Guided Feature Aligning. Finally, in the Image Decoding step, the processed CNN Visual Embeddings are decoded by a CNN decoder, upsampled, and further decoded by an additional CNN.

B. Data Preparation

In the Data Preparation step, the HR image captured by the satellite is downsampled to produce an LR image. The HR image is also used to generate a descriptive text caption. As transferring HR images from satellite to Earth is expensive and demands significant bandwidth, a Vision-Language Model (VLM) is employed onboard the satellite to generate the text caption. To minimize transmission cost and bandwidth usage, only the generated text caption and the LR image are transmitted to Earth.

C. Multimodal Encoding

In this process, the LR image is processed by both a CNN Image Encoder and a CLIP Image Encoder to generate CNN Visual Embeddings and CLIP Visual Embeddings, respectively. At the same time, the text caption is processed through a CLIP Text Encoder to produce Text Embeddings.

As discussed, remote sensing images often contain repetitive patterns and similar scene types, making it difficult for the CLIP Visual Encoder alone to effectively distinguish between them. To address this, a dual image encoder architecture is adopted to combine the semantic understanding capability of the CLIP Visual Encoder with the local feature extraction strength of the CNN. In this architecture, the CNN Visual Embeddings serve as the primary features, while the CLIP Visual Embeddings provide supplementary information to guide feature adjustment and aligning.

The image encoding process is described as follows. Let $\mathbf{L} \in \mathbb{R}^{C \times H \times W}$ represent the LR image received from the satellite. The extraction of CNN Visual Embeddings using the CNN Visual Encoder is defined by:

$$\mathbf{X}_0 = \text{VE}_{\text{CNN}}(\mathbf{L}), \quad (1)$$




Fig. 1. Overview of the proposed Text-Conditioned Remote Sensing Super-Resolution Network (Text-RSSR).

where VE_{CNN} denotes the CNN Visual Encoder, and $\mathbf{X}_0$ represents the resulting CNN Visual Embeddings.

To provide CLIP-based semantic guidance, the first token and the remaining tokens from the CLIP Visual Encoder output are separated. A Semantic Guiding (SG) process is then applied to generate a semantic guidance map for each Text-Guided Feature Aligning step in the Text-Guided Iterative Super-Resolution process [17]. The extraction of CLIP Visual Embeddings and the computation of the semantic guidance map for the i -th Text-Guided Feature Aligning step are defined as follows:

$$\mathbf{S}_i = \text{SG}_i(\mathbf{v}_0, \mathbf{v}_{\text{rm}}), \quad (2)$$

where

$$\mathbf{v}_0, \mathbf{v}_{\text{rm}} = \text{VE}_{\text{CLIP}}(\mathbf{L}).$$

Here, $\mathbf{S}_i$ and SG_i denote the CLIP Visual Embeddings and the Semantic Guiding process for the i -th Text-Guided Feature Aligning step, respectively. VE_{CLIP} denotes the CLIP Visual Encoder, $\mathbf{v}_0$ represents the first token of the resulting CLIP Visual Embeddings, and $\mathbf{v}_{\text{rm}}$ denotes the remaining tokens. Let the number of SR Units be l . Then, the number of Text-Guided Feature Aligning steps is $l + 1$, with the index $i \in [0, l]$.

The SG process is designed to modulate local feature enhancement in the Super-Resolution pipeline using semantic priors from CLIP. For each step $i \in [0, l]$, the SG process computes a semantic guidance map $\mathbf{S}_i \in \mathbb{R}^{B \times H \times W}$ for each of the l SR Units, providing explicit spatially-aware guidance aligned with the target semantics.

Let $\mathbf{v}_0 \in \mathbb{R}^{B \times C}$ denote the global semantic embedding (i.e., the [CLS] token) and $\mathbf{v}_{\text{rm}} \in \mathbb{R}^{B \times (H \times W) \times C}$ the local semantic embeddings extracted from the CLIP Visual Encoder. The SG process for the i -th step is defined as:

$$\text{SG}_i(\mathbf{v}_0, \mathbf{v}_{\text{rm}}) = \text{CS}(\Theta(\mathbf{e}_i), \Theta(\mathbf{v}_{\text{rm}})), \quad (3)$$

where $\text{CS}(\cdot, \cdot)$ denotes the cosine similarity, and $\Theta(\cdot)$ is the ℓ_2 -normalization function applied along the feature dimension. The latent embeddings $\mathbf{e}_i \in \mathbb{R}^{B \times l \times C}$ used for computing the

semantic guidance maps are derived through the following stages:

$$\begin{cases} \mathbf{p}_a = \text{MetaNet}(\mathbf{v}_0, \mathbf{p}_0), \\ [\mathbf{p}_a^*; \mathbf{p}_b^*] = \text{MHSA}([\mathbf{p}_a; \mathbf{p}_b]), \\ [\mathbf{e}_i] = \text{MLP}_e(\text{MLP}_a(\mathbf{p}_a^*) \odot \sigma(\text{MLP}_b(\mathbf{p}_b^*))). \end{cases} \quad (4)$$

Here, $\mathbf{p}_0 \in \mathbb{R}^{1 \times C}$ is a learnable global prior descriptor, and $\mathbf{p}_b \in \mathbb{R}^{(l+1) \times C}$ are learnable local unit descriptors, shared across all steps. MetaNet is an architecture to extract semantic correspondence between $\mathbf{v}_0$ and $\mathbf{p}_0$ [18]. The Multi-Head Self-Attention (MHSA) block refines the concatenated descriptors $[\mathbf{p}_a; \mathbf{p}_b]$, enabling context exchange between global and unit tokens [19]. The gating mechanism modulates the unit embeddings via element-wise product between the transformed global vector and a gated local vector. The resulting fused descriptors $\mathbf{e}_i$ are compared with the normalized local features $\mathbf{v}_{\text{rm}}$ via cosine similarity, and $i \in [0, l]$.

The output of the semantic guiding process is a semantic guidance map $\mathbf{S}_i \in \mathbb{R}^{B \times H \times W}$, where $H \times W$ corresponds to the resolution of the LR image. This map serves as a soft spatial prior for the i -th Text-Guided Feature Aligning step, guiding the model to enhance semantically meaningful regions in accordance with the text prompt. This design allows each SR Unit to adaptively incorporate global semantic context and local alignment cues, thereby improving spatial correspondence and preserving high-frequency details during the Super-Resolution reconstruction process.

The Text Embeddings are generated using the CLIP Text Encoder as follows:

$$\mathbf{T} = \text{TE}(\mathbf{t}), \quad (5)$$

where $\mathbf{t}$ denotes the caption of the HR image generated by the VLM, TE denotes Text Encoder, and $\mathbf{T}$ represents the resulting Text Embeddings.

D. Text-Guided Iterative Super-Resolution

In this step, the CNN Visual Embeddings are iteratively aligned and refined. To maintain the supervisory influence of the CLIP Visual Embeddings and Textual Embeddings throughout the entire Super-Resolution process, each iteration




Fig. 2. Overview of the Text-Guided Feature Aligning process.

begins by using these embeddings to align the CNN Visual Embeddings. The aligned CNN Visual Embeddings are then processed through an SR Unit. After the final iteration, the refined CNN Visual Embeddings undergo one final Text-Guided Feature Aligning step. Let the CNN Visual Embeddings input for the i -th iteration be $\mathbf{X}_i$, where $i \in [0, l)$. Then, the i -th iteration is defined as:

$$\mathbf{X}_{i+1} = \text{SR}(\text{TGFA}(\mathbf{X}_i)), \quad (6)$$

where TGFA denotes the Text-Guided Feature Aligning step, and SR denotes the Super-Resolution Unit.

Text-Guided Feature Aligning is designed to align and refine the CNN Visual Embeddings using both the CLIP Visual Embeddings and the CLIP Textual Embeddings. In this process, textual information provides channel-wise attention, enabling the model to identify and emphasize the most relevant feature channels for the Super-Resolution task. Following this channel-wise refinement, the semantic guidance maps are applied to perform pixel-wise refinement, offering direct high-frequency guidance to enhance spatial detail in the reconstructed image.

Let $\mathbf{T} \in \mathbb{R}^d$ be the Textual Embeddings, $\mathbf{S}_i$ the semantic guidance map for the i -th iteration, and d the dimension of CLIP's hidden state. The Text-Guided Feature Aligning process is defined as:

$$\text{TGFA}(\mathbf{X}_i) = \alpha \cdot \text{Align}(\mathbf{X}_i) + (1 - \alpha) \cdot \mathbf{X}_i, \quad (7)$$

where

$$\text{Align}(\mathbf{X}_i) = \mathbf{X}_i \odot \text{TG}(\mathbf{T}) + \mathbf{S}_i.$$

Here, α is a learnable aligning factor that controls the proportion of alignment, $\text{Align}(\mathbf{X}_i)$ denotes the aligned visual embeddings, and $\odot$ represents element-wise multiplication. The function TG is a lightweight MLP composed of a convolution layer, a ReLU activation, another convolution layer, and a Sigmoid activation, projecting the Textual Embeddings into the channel space.

The entire process of Text-Guided Feature Aligning is illustrated in Figure 2. For the Super-Resolution Unit, we adopt the Residual Token Selective Group (RTSG) from Y. Xiao *et al.*'s Top- k Token Selective Transformer (TTST) [20] architecture.

E. Image Decoding

In the Image Decoding step, the processed CNN Visual Embeddings are used to generate the super-resolved (SR)

image, which serves as the final output of the entire network. The CNN Visual Embeddings obtained from the last Text-Guided Feature Aligning are first passed through a CNN, then upsampled, and subsequently processed by another CNN to produce the SR image.

III. EXPERIMENTS

A. Implementation details

This study is implemented using PyTorch [21], and all experiments are conducted on a system equipped with an Intel Core i9-10900X CPU, 128 GiB of RAM, and a single NVIDIA GeForce RTX 4090 GPU. The operating system is Ubuntu 22.04 (Jammy) with CUDA version 12.4. The Text-RSSR framework, along with all baseline models used for comparison, adopts L1 Loss [22] as the loss function and employs the AdamW optimizer [23] with an initial learning rate of 10^{-4} . The learning rate is reduced by half at the midpoint of the training process. The downscale factor used to generate the LR image is set to 4. In this study, the training epochs were set to a moderate number to simulate a computation-constrained scenario. All models are trained from scratch, except for the CLIP components. For Text-RSSR, the Vision-Language Model (VLM) used to generate captions is LLaVa-1.5 7B [24]. The CLIP Visual Encoder and Textual Encoder used in the Multimodal Encoding step are pretrained and kept frozen during training. The initial aligning factor α in the Text-Guided Feature Aligning step is set to 0.5. The number of SR Units, l , is set to 6. The RTSG Super-Resolution Unit is initialized following the configuration provided in the original TTST paper [20]. Benchmark evaluations are performed on three public RSSR datasets: Alsat-2B dataset [25], UC Merced Land Use Dataset [26], and Aerial Image Dataset (AID) [27]. Model performance is assessed using Peak Signal-to-Noise Ratio (PSNR) [28] and Structural Similarity Index Measure (SSIM) [29]. Compared to existing methods, Text-RSSR demonstrates strong performance on RSSR tasks.

B. Dataset Description

We conducted experiments on three public RSSR datasets: Alsat-2B dataset [25], UC Merced Land Use Dataset [26], and AID [27]. According to the results, Text-RSSR demonstrates strong performance across all evaluated datasets.

The Alsat-2B dataset, introduced by A. Djerida *et al.*, includes 2,182 training samples and three testing subsets—agriculture, urban, and special scenes—containing 56,

TABLE I

QUANTITATIVE RESULTS OF THE EXPERIMENT ON ALSAT-2B DATASET. THE BEST RESULT IS HIGHLIGHTED IN **BOLD**, WHILE THE SECOND-BEST RESULT IS UNDERLINED.

Model	Parameters (M)	FLOPs (G)	Average		Agriculture		Special		Urban	
			PSNR (dB)	SSIM	PSNR (dB)	SSIM	PSNR (dB)	SSIM	PSNR (dB)	SSIM
FAT [30]	10.19	43.86	16.0099	0.3707	17.5152	0.4409	16.6609	0.3946	14.8890	0.3261
FMSR [5]	7.05	30.71	16.2615	0.3539	17.6293	0.4184	16.8002	0.3768	15.3054	0.3117
GCRDN [31]	31.83	443.80	15.8084	0.2875	17.1307	0.3931	16.3803	0.3188	14.8237	0.2258
TSFNet [32]	3.13	111.17	16.4214	<u>0.3704</u>	<u>17.7688</u>	0.4376	17.0326	0.3951	<u>15.3846</u>	0.3255
TTST [20]	18.30	76.84	16.2754	0.3542	17.6870	0.4220	16.8012	0.3764	15.3243	0.3121
Text-RSSR	22.72	76.86	<u>16.3632</u>	0.3702	18.0134	<u>0.4385</u>	<u>16.8149</u>	0.3868	15.4435	0.3346

TABLE II

QUANTITATIVE RESULTS OF THE EXPERIMENT ON UC MERCED LAND USE DATASET. THE BEST RESULT IS HIGHLIGHTED IN **BOLD**, WHILE THE SECOND-BEST RESULT IS UNDERLINED. THE UNIT FOR PSNR IS DECIBELS (DB).

Class	FAT [30]		FMSR [5]		GCRDN [31]		TSFNet [32]		TTST [20]		Text-RSSR	
	PSNR	SSIM	PSNR	SSIM	PSNR	SSIM	PSNR	SSIM	PSNR	SSIM	PSNR	SSIM
Agriculture	24.9965	0.4805	24.9226	<u>0.4757</u>	24.7018	0.4128	24.5531	0.4259	24.2153	0.4154	24.4920	0.4551
Airplane	27.8155	0.8212	28.2478	0.8320	28.2134	0.8321	28.0971	0.8279	28.4077	0.8343	28.4183	0.8350
Baseball Diamond	32.9822	0.8533	33.3289	0.8617	33.2192	0.8570	33.2576	0.8595	<u>33.4316</u>	<u>0.8626</u>	33.4700	0.8637
Beach	36.4693	0.8987	36.7805	0.9061	36.9861	0.9072	37.0140	0.9069	<u>37.0934</u>	0.9082	37.1166	0.9087
Buildings	23.6873	0.7487	24.1050	0.7656	24.2573	0.7681	23.8823	0.7552	<u>24.3830</u>	<u>0.7737</u>	24.4572	0.7772
Chaparral	25.3234	0.7543	25.6270	0.7621	25.6057	0.7607	25.6488	0.7598	25.7485	<u>0.7672</u>	<u>25.7471</u>	0.7676
Dense Residential	25.7298	0.7895	26.2259	0.8049	26.2596	0.8058	26.0135	0.7979	26.4440	0.8106	26.5231	0.8135
Forest	25.6720	0.6391	25.9337	0.6409	25.9508	0.6401	25.9675	0.6388	<u>25.9932</u>	<u>0.6451</u>	26.0015	0.6463
Freeway	27.4274	0.7525	27.7678	0.7644	27.8962	0.7647	27.5942	0.7551	<u>27.9259</u>	<u>0.7685</u>	27.9556	0.7711
Golf Course	33.5669	0.8554	33.8471	0.8625	33.9097	0.8624	33.8675	0.8617	33.9689	<u>0.8640</u>	<u>33.9649</u>	0.8643
Harbor	21.0594	0.8232	21.5372	0.8368	21.5931	0.8392	21.1637	0.8282	<u>21.8485</u>	<u>0.8495</u>	21.9199	0.8504
Intersection	25.4345	0.7533	25.8930	0.7690	25.9169	0.7680	25.7135	0.7567	<u>26.0866</u>	<u>0.7759</u>	26.1001	0.7777
Medium Residential	27.5041	0.8053	28.1087	0.8196	28.1493	0.8205	28.0127	0.8167	<u>28.3415</u>	<u>0.8242</u>	28.4243	0.8262
Mobile Home Park	17.4454	0.5973	17.6357	0.6131	17.7582	0.6147	17.4915	0.5976	<u>17.8737</u>	0.6281	17.9357	<u>0.6276</u>
Overpass	24.6758	0.6933	24.9872	0.7054	<u>25.0903</u>	0.7070	24.7842	0.6948	25.0561	0.7091	25.1763	0.7132
Parking Lot	21.1731	0.7508	21.5318	0.7631	21.5031	0.7605	21.1645	0.7452	<u>21.9376</u>	<u>0.7750</u>	21.9953	0.7788
River	28.9958	0.7852	29.4090	0.7941	29.4689	0.7959	29.4365	0.7931	<u>29.4905</u>	<u>0.7971</u>	29.5195	0.7978
Runway	30.3700	0.7884	30.5751	0.7929	30.5688	0.7899	30.4315	0.7849	<u>30.6446</u>	<u>0.7945</u>	30.8615	0.8001
Sparse Residential	27.2791	0.7770	27.5781	0.7851	27.6738	0.7873	27.5750	0.7831	27.7425	0.7904	<u>27.7201</u>	0.7891
Storage Tanks	29.3186	0.8208	29.7391	0.8316	29.7951	0.8330	29.7103	0.8287	<u>29.8461</u>	<u>0.8339</u>	29.8824	0.8358
Tennis Court	28.9392	0.8071	29.3549	0.8217	29.3263	0.8168	29.1162	0.8065	<u>29.4137</u>	<u>0.8227</u>	29.6164	0.8293
Average	26.9460	0.7617	27.2922	0.7718	27.3259	0.7688	27.1664	0.7631	<u>27.4235</u>	<u>0.7738</u>	27.4904	0.7775

282, and 239 image pairs, respectively. Each high-resolution (HR) image has a fixed resolution of 256×256 , while each low-resolution (LR) image has a resolution of 64×64 . The spatial resolution of the dataset is 2.5 meters. All models are trained for 50 epochs, and the batch size for Text-RSSR is set to 4.

The UC Merced Land Use Dataset, introduced by Y. Yang *et al.*, consists of 21 categories, each with 100 HR images. Each image has an approximate resolution of 256×256 , and a spatial resolution of 0.3 meters. For consistency in this study, all HR images are reshaped to 256×256 . We randomly select 75 images from each category for training and use the remaining 25 for testing. All models are trained for 100 epochs, and the batch size for Text-RSSR is set to 4.

The AID comprises 10,000 RGB images, each with a resolution of 600×600 pixels. Introduced by G.-S. Xia *et al.*, it contains 30 scene categories, with several hundred images per class, and a spatial resolution of 0.5 meters. In this study, 80% of the images from each category are randomly selected for training, and the remaining 20% for testing, resulting in 8,000 training images and 2,000 testing images. All HR images

are resized to 576×576 , and all LR images are resized to 144×144 . All models are trained for 20 epochs, and the batch size for Text-RSSR is set to 1.

C. Performance Comparison

To evaluate the performance of the proposed Text-RSSR model, we compared the performance of Text-RSSR with other state-of-the-art methods from 2024 and 2025, including X. Yan *et al.*'s FAT [30], Y. Xiao *et al.*'s FMSR [5], J. Sui *et al.*'s GCRDN [31], J. Wang *et al.*'s TSFNet [32] and Y. Xiao *et al.*'s TTST [20]. This study employs two evaluation metrics to assess the quality of the SR images: PSNR and SSIM.

As shown in Tables I, II, and III, Text-RSSR achieves the highest PSNR and SSIM scores across all three datasets. The learnable parameter counts and FLOPs for each model are provided in Table I. These values are computed using Python's *thop* library, with FLOPs calculated for a 64×64 RGB input image. Compared with Y. Xiao *et al.*'s TTST model, Text-RSSR improves PSNR by 0.0878 dB and SSIM by 0.0160, while increasing FLOPs by only 0.02 G. This highlights Text-RSSR's high computational efficiency. Due to the domain

TABLE III
QUANTITATIVE RESULTS OF THE EXPERIMENT ON AID. THE BEST RESULT IS HIGHLIGHTED IN **BOLD**, WHILE THE SECOND-BEST RESULT IS UNDERLINED. THE UNIT FOR PSNR IS DECIBELS (dB).

Class	FAT [30]		FMSR [5]		GCRDN [31]		TSFNet [32]		TTST [20]		Text-RSSR	
	PSNR	SSIM	PSNR	SSIM	PSNR	SSIM	PSNR	SSIM	PSNR	SSIM	PSNR	SSIM
Airport	27.7525	0.7324	<u>27.9582</u>	0.7419	27.6506	0.7277	27.7524	0.7337	27.9100	0.7384	27.9673	<u>0.7401</u>
Bare Land	34.5639	0.8269	<u>34.7455</u>	<u>0.8323</u>	34.6666	0.8289	34.6324	0.8304	<u>34.7469</u>	0.8322	34.8499	0.8328
Baseball Field	27.7234	0.7363	27.9125	0.7459	27.6471	0.7350	27.7408	0.7401	27.7956	0.7427	<u>27.9062</u>	<u>0.7454</u>
Beach	31.5206	0.7959	31.7368	0.8017	31.6311	0.7967	31.6762	0.7993	<u>31.7811</u>	<u>0.8025</u>	31.8627	0.8028
Bridge	29.5164	0.8066	<u>29.7611</u>	0.8140	29.3744	0.8033	29.5317	0.8083	29.7168	0.8122	29.7696	0.8134
Center	24.1964	0.6563	<u>24.4127</u>	0.6683	24.0441	0.6476	24.2564	0.6598	24.3907	0.6659	24.4493	0.6683
Church	23.5623	0.6633	<u>23.7468</u>	0.6747	23.4374	0.6554	23.5907	0.6661	23.7075	0.6711	23.7721	<u>0.6743</u>
Commercial	24.3682	0.6655	<u>24.5198</u>	0.6741	24.2827	0.6595	24.4117	0.6687	24.4918	0.6717	24.5252	<u>0.6737</u>
Dense Residential	25.4483	0.6691	25.5893	0.6774	25.3104	0.6597	25.4380	0.6699	25.5417	0.6724	<u>25.5836</u>	<u>0.6755</u>
Desert	35.4974	0.8603	<u>35.6769</u>	0.8646	35.5656	0.8623	35.4717	0.8622	35.5750	0.8649	35.8370	0.8652
Farmland	32.3457	0.8001	32.5888	0.8090	32.2252	0.7972	32.2969	0.8011	32.4639	0.8054	<u>32.5196</u>	<u>0.8064</u>
Forest	25.6715	0.5934	25.8384	<u>0.5994</u>	25.6358	0.5859	25.7607	0.5953	25.8171	0.5971	25.8607	0.5998
Industrial	25.3502	0.6885	<u>25.5426</u>	0.6982	25.1632	0.6782	25.3566	0.6899	25.4803	0.6942	25.5438	<u>0.6970</u>
Meadow	30.5650	0.6421	<u>30.7640</u>	0.6548	30.6964	0.6493	30.7205	0.6529	30.7427	0.6536	30.7743	<u>0.6544</u>
Medium Residential	24.1298	0.5843	24.3298	<u>0.5951</u>	24.0050	0.5725	24.2174	0.5894	24.2783	0.5911	<u>24.3296</u>	0.5952
Mountain	26.9741	0.6309	27.0518	0.6337	26.9325	0.6250	26.9786	0.6326	<u>27.0588</u>	<u>0.6346</u>	27.0870	0.6357
Park	25.2941	0.6589	25.4161	<u>0.6655</u>	25.2218	0.6533	25.3075	0.6610	25.4103	0.6638	25.4422	0.6663
Parking	24.5978	0.7433	24.9916	0.7607	24.1571	0.7263	24.6874	0.7474	24.3829	0.7408	<u>24.8294</u>	<u>0.7539</u>
Playground	30.0128	0.7827	<u>30.2509</u>	0.7918	29.7851	0.7770	30.0155	0.7848	30.1723	0.7888	30.2731	<u>0.7910</u>
Pond	27.8899	0.7152	<u>27.9814</u>	<u>0.7191</u>	27.8211	0.7123	27.8758	0.7161	27.9796	0.7186	28.0043	0.7200
Port	26.2843	0.7471	<u>26.4492</u>	0.7545	26.1971	0.7440	26.2961	0.7489	26.4259	0.7528	26.4616	<u>0.7541</u>
Railway Station	25.5554	0.6620	<u>25.7163</u>	0.6698	25.3999	0.6506	25.5058	0.6609	25.6945	0.6667	25.7389	<u>0.6695</u>
Resort	22.7878	0.6255	<u>22.9452</u>	<u>0.6342</u>	22.7181	0.6206	22.7916	0.6278	22.9095	0.6325	22.9499	0.6345
River	28.5470	0.6859	<u>28.6762</u>	<u>0.6926</u>	28.5350	0.6848	28.5804	0.6897	28.6669	0.6919	28.6964	0.6928
School	24.0299	0.6776	<u>24.2210</u>	<u>0.6871</u>	23.9095	0.6716	24.0555	0.6810	24.1896	0.6849	24.2368	0.6876
Sparse Residential	24.6960	0.5454	<u>24.8526</u>	0.5526	24.6912	0.5397	24.7662	0.5477	24.8212	0.5498	24.8562	<u>0.5524</u>
Square	28.3506	0.7540	<u>28.5661</u>	0.7636	28.1843	0.7483	28.3606	0.7565	28.5292	0.7613	28.5974	<u>0.7634</u>
Stadium	26.0934	0.7377	<u>26.3204</u>	<u>0.7486</u>	25.9347	0.7310	26.1205	0.7399	26.2661	0.7471	26.3626	0.7502
Storage Tanks	24.8955	0.6788	25.0844	0.6894	24.7784	0.6727	24.9263	0.6815	25.0396	0.6863	<u>25.0742</u>	<u>0.6881</u>
Viaduct	25.6030	0.6395	<u>25.7718</u>	0.6497	25.4949	0.6314	25.6565	0.6427	25.7504	0.6467	25.7897	0.6497
Average	27.2062	0.7034	<u>27.3927</u>	0.7120	27.1123	0.6981	27.2373	0.7060	27.3347	0.7092	27.4089	<u>0.7116</u>

shift between LR and HR images in the Alsat-2B dataset, X. Yan *et al.*'s FAT and J. Sui *et al.*'s GCRDN exhibit a substantial drop in performance (PSNR <16 dB), indicating their limited ability to handle domain shift. Moreover, in the UC Merced Land Use Dataset and AID, which include a greater variety of image categories, the benefits of text guidance are more pronounced. In the UC Merced Land Use Dataset, containing 21 image categories, Text-RSSR achieves a PSNR of 27.4904 dB and an SSIM of 0.7775, surpassing the second-best results by 0.669 dB and 0.0037, respectively. In AID, which comprises 30 image categories, Text-RSSR attains a PSNR of 27.4089 dB and an SSIM of 0.7116, exceeding the second-best results by 0.0742 dB and 0.0024, respectively. These quantitative findings underscore the effectiveness of incorporating text guidance.

To further evaluate the visual quality of the SR images, we compare the results of Text-RSSR with those of other models across all three datasets, as shown in Figure 3. In addition to achieving higher PSNR and SSIM scores, Text-RSSR also demonstrates superior visual fidelity. When a domain shift exists between LR and HR images, as illustrated in the first and second rows, Text-RSSR successfully restores the correct color domain of the HR images. In cases involving repetitive line patterns, as seen in the fourth and sixth rows, Text-RSSR reconstructs the lines directly rather than simply emphasizing them. Moreover, Text-RSSR performs particularly well on

parking lot images, achieving state-of-the-art results compared with other methods, as shown in the last two rows.

TABLE IV
QUANTITATIVE RESULTS OF THE ABLATION STUDY ON TEXT-GUIDED FEATURE ALIGNING FOR THE UC MERCED LAND USE DATASET. CNN V. E. REFERS TO CNN VISUAL EMBEDDINGS, T. E. TO TEXTUAL EMBEDDINGS, *Hierarchy* TO THE HIERARCHICAL ARCHITECTURE OF THE TEXT-GUIDED FEATURE ALIGNING PROCESS, AND A. F. TO THE LEARNABLE ALIGNING FACTOR. THE BEST RESULT IS SHOWN IN **BOLD**, WHILE THE SECOND-BEST RESULT IS UNDERLINED.

Model Settings				Metrics	
CNN V. E.	T. E.	Hierarchy	A. F.	PSNR (dB)	SSIM
✓				27.4464	0.7726
	✓			27.4506	0.7744
✓	✓			27.4580	0.7750
✓	✓	✓		<u>27.4771</u>	<u>0.7773</u>
✓	✓	✓	✓	27.4904	0.7775

D. Ablation Study

To further evaluate the effectiveness of each components of Text-Guided Feature Aligning process, this study conducts an ablation study on Text-Guided Feature Aligning on the UC Merced Land Use Dataset. We first construct a model with a single Text-Guided Feature Aligning process applied before




Fig. 3. Qualitative results of the experiment on all three datasets.

the first SR Unit, using only CNN Visual Embeddings and removing the Textual Embeddings. Next, we build a second model that uses only Textual Embeddings, excluding the CNN Visual Embeddings. We then create a third model that incorporates both CNN Visual Embeddings and Textual Embeddings. Subsequently, we design a hierarchical architecture by adding a Text-Guided Feature Aligning process after each SR Unit. Finally, we introduce a learnable aligning factor for each Text-Guided Feature Aligning process to fine-tune the alignment, resulting in the complete Text-RSSR model.

The quantitative results of the ablation study are presented in Table IV. As shown, incorporating CNN Visual Embeddings as pixel-wise guidance or Textual Embeddings as channel-wise guidance both enhance performance, with Textual Embeddings yielding a slightly greater improvement (0.0042 dB in PSNR and 0.0018 in SSIM) compared to CNN Visual Embeddings. Combining both types of embeddings further boosts reconstruction quality, demonstrating that pixel-wise and channel-wise guidance provide complementary information for the feature aligning process. Introducing the hierarchical architecture

brings an additional performance gain (0.0191 dB in PSNR and 0.0023 in SSIM), indicating that multi-level alignment enhances the integration of visual and textual cues. Finally, adding a learnable aligning factor provides the Text-RSSR model with greater flexibility in leveraging text guidance, leading to a further improvement in performance.

IV. CONCLUSION

This study proposes Text-RSSR, a novel framework for leveraging text guidance in remote sensing image super-resolution. By integrating multimodal cues through CLIP embeddings and iteratively aligning visual and textual features, Text-RSSR effectively enhances the reconstruction of fine-grained spatial details. The four-step pipeline (Data Preparation, Multimodal Encoding, Text-Guided Iterative Super-Resolution, and Image Decoding) demonstrates how semantic information from text captions can be systematically incorporated into the super-resolution process. This approach highlights the potential of multimodal learning to bridge the gap between low-resolution and high-resolution imagery, offering a promising direction for advancing remote sensing applications that demand both high fidelity and semantic consistency.

The experimental results further validate the effectiveness of the proposed Text-RSSR framework. Across all three benchmark datasets: Alsat-2B, UC Merced Land Use, and AID, Text-RSSR consistently outperforms state-of-the-art methods in both PSNR and SSIM, while maintaining competitive computational efficiency. Notably, Text-RSSR achieves substantial improvements in scenarios with pronounced domain shifts and complex scene variations, underscoring the robustness of its text-guided alignment mechanism. In the UC Merced and AID datasets, which cover diverse scene categories, Text-RSSR demonstrates particularly strong gains, highlighting the value of multimodal guidance in handling heterogeneity in remote sensing imagery. Moreover, qualitative evaluations reveal that Text-RSSR excels in restoring fine structural details and semantic consistency, such as reconstructing repetitive line patterns and preserving correct color domains under challenging conditions. The ablation study further confirms that the hierarchical integration of CNN Visual and Textual Embeddings, together with a learnable aligning factor, is critical to achieving optimal performance. Collectively, these findings establish Text-RSSR as a highly effective and efficient framework for advancing the state of remote sensing image super-resolution.

REFERENCES

- [1] R. Keys, "Cubic convolution interpolation for digital image processing," *IEEE Transactions on Acoustics, Speech, and Signal Processing*, vol. 29, no. 6, pp. 1153–1160, 1981.
- [2] J. Yang, J. Wright, T. S. Huang, and L. Yu, "Image super-resolution via sparse representation," *IEEE Transactions on Image Processing*, vol. 19, no. 11, pp. 2861–2873, 2010.
- [3] C. Dong, C. C. Loy, K. He, and X. Tang, "Image super-resolution using deep convolutional networks," *arXiv preprint arXiv:1501.00092*, pp. 1–14, 2015.
- [4] B. Lim, S. Son, H. Kim, S. Nah, and K. M. Lee, "Enhanced deep residual networks for single image super-resolution," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR) Workshops*, 2017, pp. 1132–1140.
- [5] Y. Xiao, Q. Yuan, K. Jiang, Y. Chen, Q. Zhang, and C.-W. Lin, "Frequency-assisted mamba for remote sensing image super-resolution," *IEEE Transactions on Multimedia*, vol. 26, pp. 1–14, 2024.
- [6] T. Wu, R. Zhao, M. Lv, Z. Jia, L. Li, and M. Liu, "Efficient mamba-attention network for remote sensing image super-resolution," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 63, pp. 1–14, 2025.
- [7] C. Ledig, L. Theis, F. Huszar, J. Caballero, A. Cunningham, A. Acosta, A. Aitken, A. Tejani, J. Totz, Z. Wang, and W. Shi, "Photo-realistic single image super-resolution using a generative adversarial network," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017, pp. 4681–4690.
- [8] S. Lei, Z. Shi, and Z. Zou, "Coupled adversarial training for remote sensing image super-resolution," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 58, no. 5, pp. 3633–3643, 2020.
- [9] H. Li, Y. Yang, M. Chang, H. Feng, Z. Xu, Q. Li, and Y. Chen, "Srdiff: Single image super-resolution with diffusion probabilistic models," *arXiv preprint arXiv:2104.14951*, pp. 1–15, 2021.
- [10] B. Kavar, M. Elad, S. Ermon, and J. Song, "Denoising diffusion restoration models," *arXiv preprint arXiv:2201.11793*, pp. 1–20, 2022.
- [11] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, "Learning transferable visual models from natural language supervision," *arXiv preprint*, 2021.
- [12] Y. Yuan, X. Meng, W. Sun, G. Yang, L. Wang, J. Peng, and Y. Wang, "Multi-resolution collaborative fusion of sar, multispectral and hyperspectral images for coastal wetlands mapping," *Remote Sensing*, vol. 14, no. 14, p. 3492, 2022.
- [13] J. F. Hu, T. Z. Huang, and L. J. Deng, "Fusformer: A transformer-based fusion approach for hyperspectral image super-resolution," *arXiv preprint arXiv:2109.02079*, pp. 1–15, 2021.
- [14] U. Jain, A. Wilson, and V. Gulshan, "Multimodal contrastive learning for remote sensing tasks," in *Proceedings of the 36th Conference on Neural Information Processing Systems (NeurIPS)*. NeurIPS, 2022, pp. 1–12.
- [15] A. Fuller, K. Millard, and J. R. Green, "Croma: Remote sensing representations with contrastive radar-optical masked autoencoders," in *Proceedings of the 37th Conference on Neural Information Processing Systems (NeurIPS)*. NeurIPS, 2023, pp. 1–12.
- [16] K. V. Gandikota and P. Chandramouli, "Text-guided explorable image super-resolution," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. IEEE, 2024, pp. 1–10.
- [17] B. Chen, K. Chen, M. Yang, Z. Zou, and Z. Shi, "Segsr: Integrating semantic knowledge into remote sensing image super-resolution via vision-language model."
- [18] S. Ravi and H. Larochelle, "Optimization as a model for few-shot learning," in *International Conference on Learning Representations (ICLR)*, 2017.
- [19] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. Gomez, Kaiser, and I. Polosukhin, "Attention is all you need," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017, pp. 5998–6008.
- [20] Y. Xiao, Q. Yuan, K. Jiang, J. He, C. W. Lin, and L. Zhang, "Tst: A top-k token selective transformer for efficient self-attention," *IEEE Transactions on Image Processing*, vol. 33, pp. 738–752, 2024.
- [21] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, A. Desmaison, A. Köpf, E. Yang, Z. DeVito, M. Raison, A. Tejani, S. Chilamkurthy, B. Steiner, L. Fang, J. Bai, and S. Chintala, "Pytorch: An imperative style, high-performance deep learning library,"

- 2019.
- [22] H. Zhao, O. Gallo, I. Frosio, and J. Kautz, “Loss functions for image restoration with neural networks,” *IEEE Transactions on Computational Imaging*, vol. 3, no. 1, pp. 47–57, 2017.
 - [23] D. Kingma and J. Ba, “Adam: A method for stochastic optimization,” *International Conference on Learning Representations (ICLR)*, 2015.
 - [24] H. Liu, C. Li, Y. Li, and Y. Lee, “Improved baselines with visual instruction tuning (llava-1.5),” *arXiv preprint*, 2024.
 - [25] A. Djerida, K. Djerriri, M. Karoui, and M. Larabi, “A new public alsat-2b dataset for single-image super-resolution.”
 - [26] Y. Yang and S. Newsam, “Bag-of-visual-words and spatial extensions for land-use classification,” in *ACM SIGSPATIAL International Conference on Advances in Geographic Information Systems (ACM GIS)*, 2010.
 - [27] G.-S. Xia, J. Hu, F. Hu, B. Shi, X. Bai, Y. Zhong, L. Zhang, and X. Lu, “Aid: A benchmark data set for performance evaluation of aerial scene classification,” *IEEE Transactions on Geoscience and Remote Sensing*, vol. 55, no. 7, pp. 3965–3981, 2017.
 - [28] A. N. Netravali and B. G. Haskell, “Digital pictures: Representation and compression,” in *Springer*, 1988.
 - [29] Z. Wang, A. Bovik, H. Sheikh, and E. Simoncelli, “Image quality assessment: From error visibility to structural similarity,” *IEEE Transactions on Image Processing*, vol. 13, no. 4, pp. 600–612, 2004.
 - [30] X. Yan, J. Chen, Q. Xu, and W. Li, “Mitigating texture bias: A remote sensing super-resolution method focusing on high-frequency texture reconstruction,” *IEEE Transactions on Geoscience and Remote Sensing*, 2025.
 - [31] J. Sui, X. Ma, X. Zhang, and M. O. Pun, “Global context-driven residual dense network for remote sensing image super-resolution,” *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, vol. 16, pp. 4457–4468, 2023.
 - [32] J. Wang, Y. Lu, S. Wang, B. Wang, X. Wang, and T. Long, “Two-stage spatial-frequency joint learning for large-factor remote sensing image super-resolution,” *IEEE Transactions on Geoscience and Remote Sensing*, vol. 62, pp. 1–14, 2024.